



```
In [1]: import pandas as pd
!pip install imblearn
```

```
Requirement already satisfied: imblearn in /opt/anaconda3/lib/python3.12/site-p
ackages (0.0)
Requirement already satisfied: imbalanced-learn in /opt/anaconda3/lib/python3.1
2/site-packages (from imblearn) (0.12.3)
Requirement already satisfied: numpy>=1.17.3 in /opt/anaconda3/lib/python3.12/s
ite-packages (from imbalanced-learn->imblearn) (1.26.4)
Requirement already satisfied: scipy>=1.5.0 in /opt/anaconda3/lib/python3.12/si
te-packages (from imbalanced-learn->imblearn) (1.13.1)
Requirement already satisfied: scikit-learn>=1.0.2 in /opt/anaconda3/lib/python
3.12/site-packages (from imbalanced-learn->imblearn) (1.5.1)
Requirement already satisfied: joblib>=1.1.1 in /opt/anaconda3/lib/python3.12/s
ite-packages (from imbalanced-learn->imblearn) (1.4.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/anaconda3/lib/pytho
n3.12/site-packages (from imbalanced-learn->imblearn) (3.5.0)
```

```
In [2]: heart=pd.read_csv("/Users/ronak/Library/Containers/com.microsoft.Excel/Data/Dc
```

```
In [3]: heart.head()
```

```
Out[3]:
```

	General_Health	Checkup	Exercise	Heart_Disease	Skin_Cancer	Other_Cance
0	Poor	Within the past 2 years	No	No	No	N
1	Very Good	Within the past year	No	Yes	No	N
2	Very Good	Within the past year	Yes	No	No	N
3	Poor	Within the past year	Yes	Yes	No	N
4	Good	Within the past year	No	No	No	N

```
In [4]: heart.shape
```

```
Out[4]: (308854, 19)
```

```
In [5]: heart.isnull().sum()
```

```
Out[5]: General_Health      0
        Checkup            0
        Exercise           0
        Heart_Disease      0
        Skin_Cancer        0
        Other_Cancer       0
        Depression         0
        Diabetes           0
        Arthritis          0
        Sex                0
        Age_Category       0
        Height_(cm)        0
        Weight_(kg)        0
        BMI                0
        Smoking_History    0
        Alcohol_Consumption 0
        Fruit_Consumption  0
        Green_Vegetables_Consumption 0
        FriedPotato_Consumption 0
        dtype: int64
```

## Data Preprocessing

```
In [6]: # Converging the column names into lower case and replacing the space with an underscore
heart.columns = heart.columns.str.lower().str.replace(" ", "_")

#Changing the name of a big column

heart.rename(columns = {'height_(cm)' : 'height', 'weight_(kg)' : 'weight', 'green_vegetables_consumption' : 'green_vegetables_consumption'})
```

```
In [7]: heart.head()
```

```
Out[7]:
```

	general_health	checkup	exercise	heart_disease	skin_cancer	other_cancer
0	Poor	Within the past 2 years	No	No	No	No
1	Very Good	Within the past year	No	Yes	No	No
2	Very Good	Within the past year	Yes	No	No	No
3	Poor	Within the past year	Yes	Yes	No	No
4	Good	Within the past year	No	No	No	No

```
In [8]: heart['checkup'] = heart['checkup'].replace('Within the past 2 years', 'Past 2 years')
heart['checkup'] = heart['checkup'].replace('Within the past year', 'Past 1 year')
heart['checkup'] = heart['checkup'].replace('Within the past 5 years', 'Past 5 years')
heart['checkup'] = heart['checkup'].replace('5 or more years ago', 'More than 5 years ago')

heart['diabetes'] = heart['diabetes'].replace('No, pre-diabetes or borderline diabetes', 'No')
heart['diabetes'] = heart['diabetes'].replace('Yes, but female told only during pregnancy', 'Yes')

heart['age_category'] = heart['age_category'].replace('18-24', 'Young')
heart['age_category'] = heart['age_category'].replace('25-29', 'Adult')
heart['age_category'] = heart['age_category'].replace('30-34', 'Adult')
heart['age_category'] = heart['age_category'].replace('35-39', 'Adult')
heart['age_category'] = heart['age_category'].replace('40-44', 'Mid-Aged')
heart['age_category'] = heart['age_category'].replace('45-49', 'Mid-Aged')
heart['age_category'] = heart['age_category'].replace('50-54', 'Mid-Aged')
heart['age_category'] = heart['age_category'].replace('55-59', 'Senior-Adult')
heart['age_category'] = heart['age_category'].replace('60-64', 'Senior-Adult')
heart['age_category'] = heart['age_category'].replace('65-69', 'Elderly')
heart['age_category'] = heart['age_category'].replace('70-74', 'Elderly')
heart['age_category'] = heart['age_category'].replace('75-79', 'Elderly')
heart['age_category'] = heart['age_category'].replace('80+', 'Elderly')
```

```
In [9]: heart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 308854 entries, 0 to 308853
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   general_health         308854 non-null object
1   checkup                308854 non-null object
2   exercise               308854 non-null object
3   heart_disease          308854 non-null object
4   skin_cancer            308854 non-null object
5   other_cancer           308854 non-null object
6   depression              308854 non-null object
7   diabetes               308854 non-null object
8   arthritis              308854 non-null object
9   sex                   308854 non-null object
10  age_category           308854 non-null object
11  height                 308854 non-null float64
12  weight                 308854 non-null float64
13  bmi                   308854 non-null float64
14  smoking_history        308854 non-null object
15  alcohol_consumption     308854 non-null float64
16  fruit_consumption       308854 non-null float64
17  vegetables_consumption  308854 non-null float64
18  potato_consumption      308854 non-null float64
dtypes: float64(7), object(12)
memory usage: 44.8+ MB
```

```
In [10]: # Visualization
import plotly.express as px
```

```
import plotly.subplots as sp
import plotly.graph_objs as go
import matplotlib.pyplot as plt
colors = px.colors.sequential.Plasma_r
```

```
In [11]: fig1 = px.histogram(heart, x="general_health", color = 'general_health', color_
fig1.update_layout(plot_bgcolor='white')
fig1.show()
print('\n', "="*90, '\n')

fig2 = px.histogram(heart, x="general_health", color = 'heart_disease', color_
fig2.update_layout(plot_bgcolor='white')
fig2.show()
print('\n', "="*90, '\n')
```

```
=====
=====
```

```
=====
=====
```

```
In [12]: age_category_counts = heart['age_category'].value_counts()
fig_1 = px.bar(x=age_category_counts.index, y=age_category_counts.values, color=age_category_counts.index)
fig_1.update_layout(title="1. Distribution of Age Categories in the Dataset", x=age_category_counts.index)
fig_1.show()
print('\n', "="*80, '\n')

fig_2 = px.histogram(heart, x="age_category", color='heart_disease', barmode='group')
fig_2.update_layout(xaxis_title="age_category", yaxis_title="Count", legend_title="Heart Disease")
fig_2.show()
print('\n', "="*80, '\n')

grouped_data = heart.groupby(['age_category', 'heart_disease'], as_index=False)
fig_3 = px.bar(grouped_data, x='age_category', y='bmi', color='heart_disease', barmode='group')
fig_3.update_layout(xaxis_title="Age Group", yaxis_title="Average BMI", legend_title="Heart Disease")
fig_3.show()
```

=====

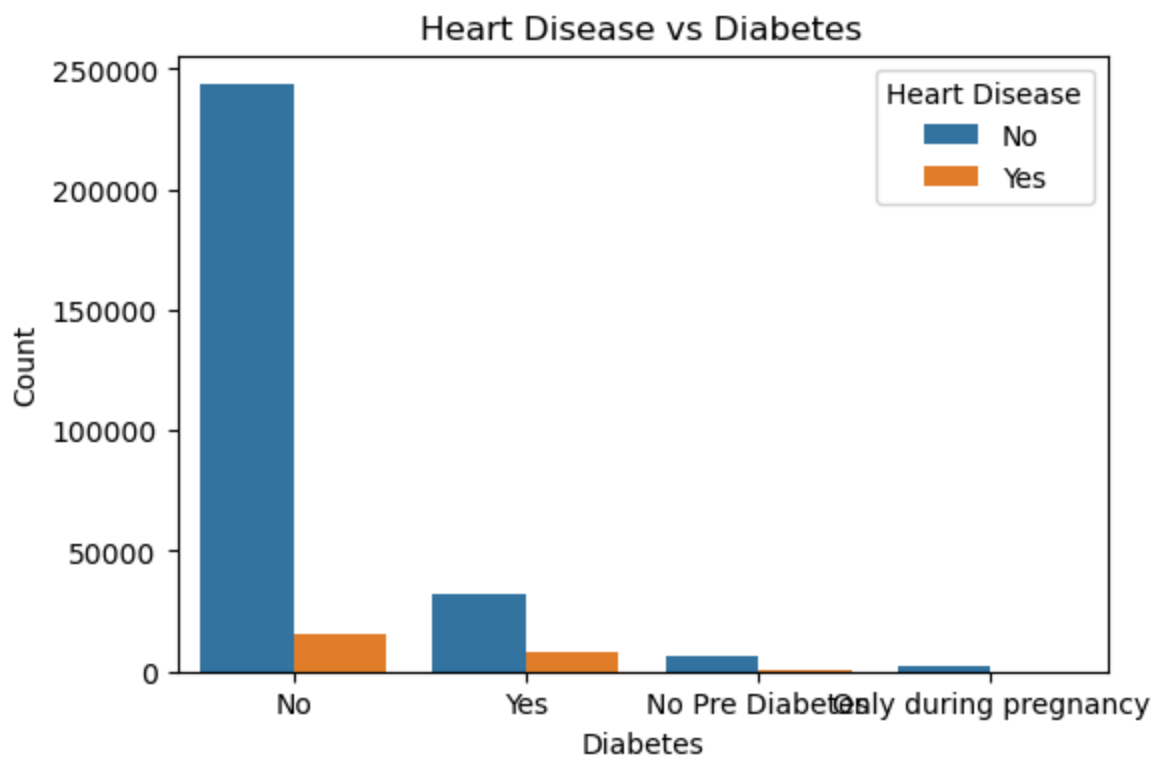
=

=====

=

```
In [13]: import seaborn as sns
plt.figure(figsize=(6,4))
sns.countplot(data=heart, x="diabetes", hue="heart_disease")
plt.title("Heart Disease vs Diabetes")
plt.xlabel("Diabetes")
plt.ylabel("Count")
plt.legend(title="Heart Disease", labels=["No", "Yes"])
plt.show()
```





In [ ]:

In [ ]:

In [ ]:

```
In [14]: col = ['alcohol_consumption', 'fruit_consumption', 'vegetables_consumption', '
for i in col:
    heart[i] = heart[i].astype(int)
```

```
In [15]: # Define BMI ranges and labels for each group
bmi_bins = [12.02, 18.3, 26.85, 31.58, 37.8, 100]
bmi_labels = ['Underweight', 'Normal weight', 'Overweight', 'Obese I', 'Obese
heart['bmi_group'] = pd.cut(heart['bmi'], bins=bmi_bins, labels=bmi_labels, ri
```

```
In [16]: column_to_move = heart.pop('bmi_group')
heart.insert(14, 'bmi_group', column_to_move)
```

```
In [17]: heart['bmi_group'] = heart['bmi_group'].astype('object')
```

```
In [18]: # Import the OneHotEncoder class from scikit-learn
from sklearn.preprocessing import OneHotEncoder
heart['heart_disease'] = heart['heart_disease'].map({'Yes':1, 'No':0})
cat=['sex', 'smoking_history']

OH_Encoder = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
```

```
OH = OH_Encoder.fit_transform(heart[cat])
cols = OH_Encoder.get_feature_names_out(cat)
OH = pd.DataFrame(OH, columns=cols)
heart = heart.drop(cat,axis=1)
heart = pd.concat([heart, OH], axis =1)
```

```
In [19]: from sklearn.preprocessing import LabelEncoder
categorical_columns = ['general_health', 'checkup', 'exercise', 'skin_cancer',

# Initialize LabelEncoder

label_encoder = LabelEncoder()

# Apply label encoding to each ordinal categorical column

for col in categorical_columns:
    heart[col] = label_encoder.fit_transform(heart[col])
```

```
In [20]: heart.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 308854 entries, 0 to 308853
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   general_health                        308854 non-null  int64
1   checkup                              308854 non-null  int64
2   exercise                             308854 non-null  int64
3   heart_disease                        308854 non-null  int64
4   skin_cancer                          308854 non-null  int64
5   other_cancer                         308854 non-null  int64
6   depression                           308854 non-null  int64
7   diabetes                             308854 non-null  int64
8   arthritis                            308854 non-null  int64
9   age_category                         308854 non-null  int64
10  height                               308854 non-null  float64
11  weight                               308854 non-null  float64
12  bmi                                  308854 non-null  float64
13  bmi_group                            308854 non-null  int64
14  alcohol_consumption                  308854 non-null  int64
15  fruit_consumption                    308854 non-null  int64
16  vegetables_consumption                308854 non-null  int64
17  potato_consumption                   308854 non-null  int64
18  sex_Female                           308854 non-null  float64
19  sex_Male                             308854 non-null  float64
20  smoking_history_No                   308854 non-null  float64
21  smoking_history_Yes                  308854 non-null  float64
dtypes: float64(7), int64(15)
memory usage: 51.8 MB
```

```
In [21]: # Compute correlation only for numerical features
corr = heart.corr(numeric_only=True)
```

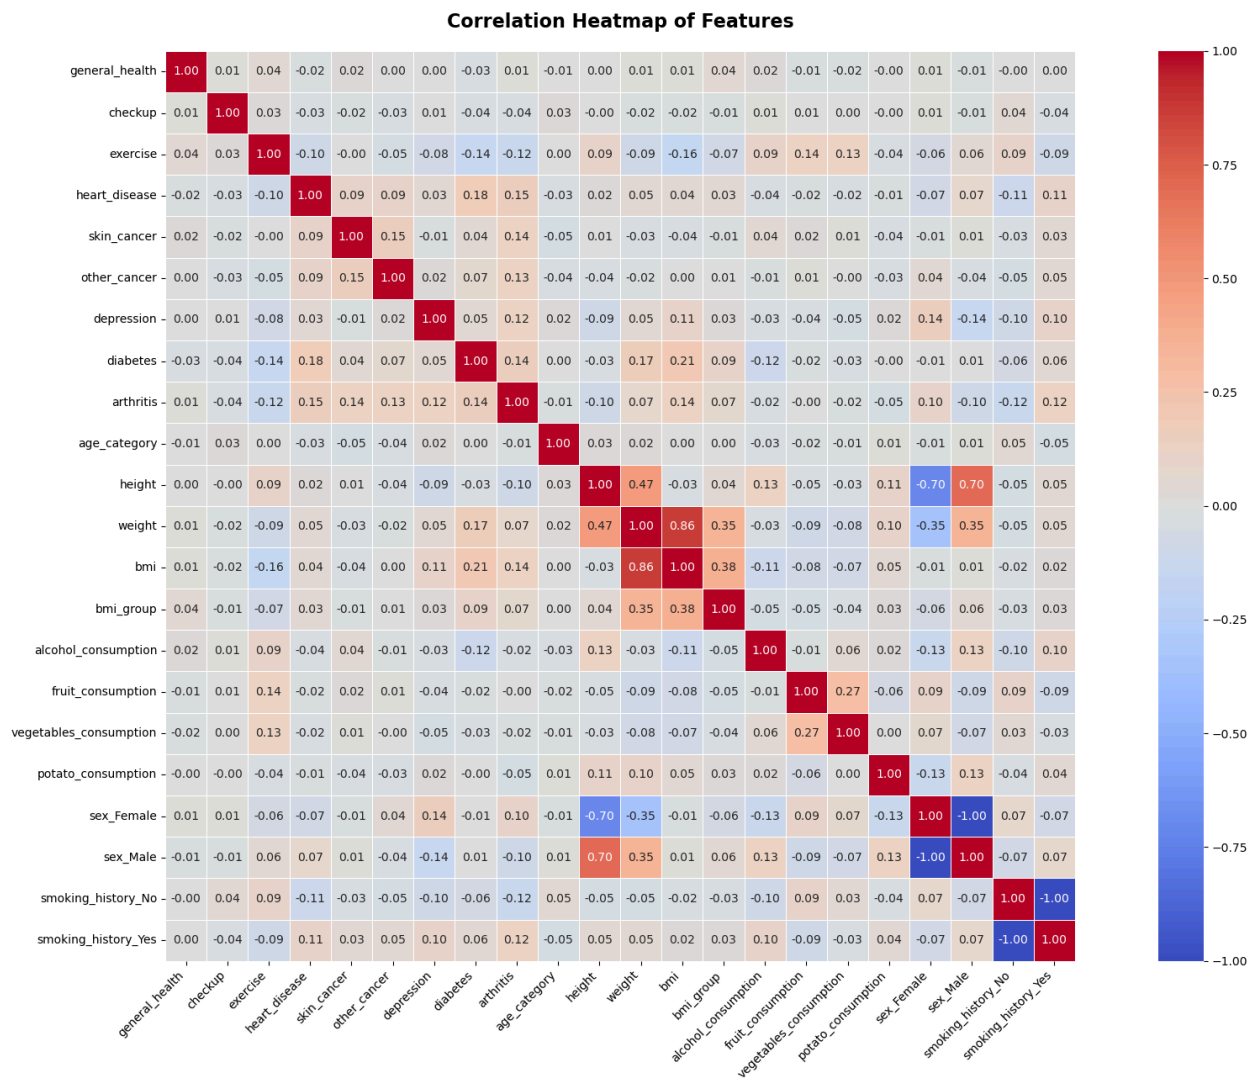
```

# Set figure size
plt.figure(figsize=(20,13))

# Draw heatmap
sns.heatmap(
    corr,
    annot=True,           # Show correlation values
    fmt=".2f",            # Format decimals
    cmap="coolwarm",      # Color scheme
    cbar=True,            # Show color bar
    square=True,          # Make cells square
    linewidths=0.5        # Add cell borders
)

plt.title("Correlation Heatmap of Features", fontsize=16, fontweight="bold", p
plt.xticks(rotation=45, ha="right")
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()

```



In [22]: heart["heart\_disease"].value\_counts()

```
Out[22]: heart_disease
0      283883
1       24971
Name: count, dtype: int64
```

```
In [23]: X = heart.drop("heart_disease", axis = 1)
y = heart['heart_disease']
```

```
In [24]: from imblearn.over_sampling import SMOTE
smote = SMOTE(random_state=42)
X_balanced, y_balanced = smote.fit_resample(X, y)
```

```
In [25]: from collections import Counter
print("Before SMOTE:", Counter(y))

# After SMOTE
print("After SMOTE:", Counter(y_balanced))

# Convert to DataFrame for better visualization
before = pd.Series(y).value_counts()
after = pd.Series(y_balanced).value_counts()

print("\nClass distribution before SMOTE:\n", before)
print("\nClass distribution after SMOTE:\n", after)
```

```
Before SMOTE: Counter({0: 283883, 1: 24971})
```

```
After SMOTE: Counter({0: 283883, 1: 283883})
```

```
Class distribution before SMOTE:
```

```
heart_disease
0      283883
1       24971
Name: count, dtype: int64
```

```
Class distribution after SMOTE:
```

```
heart_disease
0      283883
1      283883
Name: count, dtype: int64
```

```
In [26]: import matplotlib.pyplot as plt

fig, axes = plt.subplots(1, 2, figsize=(10,4))

# Before SMOTE
before.plot(kind="bar", ax=axes[0], color=["skyblue", "salmon"])
axes[0].set_title("Before SMOTE")
axes[0].set_xlabel("Class")
axes[0].set_ylabel("Count")


# After SMOTE
after.plot(kind="bar", ax=axes[1], color=["skyblue", "salmon"])
axes[1].set_title("After SMOTE")
```

```

axes[1].set_xlabel("Class")
axes[1].set_ylabel("Count")

plt.tight_layout()
plt.show()

```



```

In [27]: from sklearn.model_selection import train_test_split
# Splitting the data into training and testing sets for diabetes balanced

X_train, X_test, y_train, y_test = train_test_split(X_balanced, y_balanced, te

```

```

In [28]: from sklearn.preprocessing import StandardScaler
scaler_d = StandardScaler()
X_train_scaled = scaler_d.fit_transform(X_train)
X_test_scaled = scaler_d.transform(X_test)

```

```

In [29]: from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, f1_score, roc_auc_score, classification_report
from xgboost import XGBClassifier

# Define models in a dictionary
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(),
    "XGBoost": XGBClassifier(use_label_encoder=False, eval_metric='logloss', r

}

```

```

In [30]: # Dictionary to store results
results = {}

# Train and evaluate each model

```

```

for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    # Calculate metrics
    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    auc = roc_auc_score(y_test, model.predict_proba(X_test_scaled)[:,-1])

    results[name] = {"Accuracy": acc, "F1-score": f1, "ROC-AUC": auc}

    print(f"=== {name} ===")
    print(classification_report(y_test, y_pred))
    print("\n")

```

```

=== Logistic Regression ===

```

	precision	recall	f1-score	support
0	0.72	0.72	0.72	85071
1	0.72	0.72	0.72	85259
accuracy			0.72	170330
macro avg	0.72	0.72	0.72	170330
weighted avg	0.72	0.72	0.72	170330

```

=== Random Forest ===

```

	precision	recall	f1-score	support
0	0.92	0.95	0.94	85071
1	0.95	0.92	0.93	85259
accuracy			0.93	170330
macro avg	0.94	0.93	0.93	170330
weighted avg	0.94	0.93	0.93	170330

```

=== Gradient Boosting ===

```

	precision	recall	f1-score	support
0	0.87	0.89	0.88	85071
1	0.89	0.86	0.88	85259
accuracy			0.88	170330
macro avg	0.88	0.88	0.88	170330
weighted avg	0.88	0.88	0.88	170330

```
/opt/anaconda3/lib/python3.12/site-packages/xgboost/training.py:183: UserWarning:
```

```
[19:19:59] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
=== XGBoost ===
```

	precision	recall	f1-score	support
0	0.90	0.94	0.92	85071
1	0.94	0.90	0.92	85259
accuracy			0.92	170330
macro avg	0.92	0.92	0.92	170330
weighted avg	0.92	0.92	0.92	170330

```
In [32]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Store results
results = []

# Train & evaluate sklearn models
for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)
    y_pred_prob = model.predict_proba(X_test_scaled)[:, 1]

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    auc = roc_auc_score(y_test, y_pred_prob)

    results.append({"Model": name, "Accuracy": acc, "F1-score": f1, "ROC-AUC": auc})

# Convert results to DataFrame
df_results = pd.DataFrame(results)
print(df_results)

# === Visualization ===
plt.figure(figsize=(10,6))
df_melted = df_results.melt(id_vars="Model", var_name="Metric", value_name="Score")


sns.barplot(data=df_melted, x="Model", y="Score", hue="Metric")
plt.title("Model Performance Comparison")
plt.xticks(rotation=30)
plt.ylim(0,1)
```

```
plt.show()
```

```
/opt/anaconda3/lib/python3.12/site-packages/xgboost/training.py:183: UserWarning:
```

```
[19:27:50] WARNING: /Users/runner/work/xgboost/xgboost/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

	Model	Accuracy	F1-score	ROC-AUC
0	Logistic Regression	0.721893	0.721821	0.800895
1	Random Forest	0.934873	0.934060	0.981420
2	Gradient Boosting	0.876939	0.875495	0.952148
3	XGBoost	0.919832	0.917980	0.974887



In [ ]: